

Shijia Yang

shijiay@stanford.edu • (626) 861-1312 • Menlo Park, California

EDUCATION

Stanford University

May 2023 - Present

- Master of Science in Computer Science with GPA: N/A

University of California, Berkeley

September 2019 - May 2023

- Bachelor of Art in Computer Science with GPA: 3.9/4.0
- Relevant Coursework:
 - Applied Theory - 3D Vision (Graduate) [A+], Machine Learning, DNN, Computer Graphics
 - Theory - Optimization, Linear Algebra, Probability, Discrete Mathematics, Multivariate Calculus

RESEARCH EXPERIENCE

Stanford Vision and Learning Lab

April 2023 - Present

Graduate Researcher, advised by Prof. Jiajun Wu

Stanford, CA

- Researching on 3D robotic applications with large language models

Mechanical Systems Control Lab

March 2021 - May 2023

Undergraduate Researcher, advised by Prof. Masayoshi Tomizuka

Berkeley, CA

- Researched on multi-modal multi-task learning for autonomous vehicles
- Investigated implicit and explicit neural representation with long-term temporal images for improved 3D detection

Berkeley Artificial Intelligence Lab

March 2021 - May 2023

Undergraduate Researcher, advised by Prof. Kurt Keutzer & Prof. Joseph E. Gonzalez

Berkeley, CA

- Researched on data and training efficiency for 3D vision and efficient prompt tuning of language-vision models
- Conducting semi-supervised and self-supervised experiments for point-cloud classification, segmentation

PUBLICATIONS

[1] Image2Point: 3D Point-Cloud Understanding with 2D Image Pretrained Models

*Shijia Yang**, *Chenfeng Xu**, *Tomer Galanti*, *Bichen Wu*, *Xiangyu Yue*, *Bohan Zhai*, *Wei Zhan*, *Peter Vajda*, *Kurt Keutzer*, *Masayoshi Tomizuka*
European Conference on Computer Vision (ECCV) 2022

- Proposed a simple pretraining scheme of using 2D model weights for 3D point-cloud understanding
- Brought consistent improvement on various baselines, including ViT and PointNet++, and 10.0% improvement for few-shot classification on the ModelNet40 dataset
- Explained feasibility of image-point-cloud transfer from the aspect of neural collapse

[2] Time Will Tell: New Outlooks and a Baseline for Temporal Multi-View 3D Object Detection

Jinhyung Park*, Chenfeng Xu*, **Shijia Yang**, Kurt Keutzer, Kris Kitani, Masayoshi Tomizuka, Wei Zhan

International Conference on Learning Representations (ICLR) 2023 Top-5%

- Formulated temporal camera-only 3D detection as multi-view stereo matching problem and observed problems of previous methods regarding the low matching resolution and limited history usage
- Proposed SOLOFusion which generate a cost volume with a long history of image observations with coarse but efficient matching resolution, then augment with short-term, fine-grained matching resolution
- SOLOFusion sets new state of-the-art on nuScenes, achieving first place on the test set and outperforming previous best art by 5.2% mAP and 3.7% NDS on the validation set

[3] Multitask Vision-Language Prompt Tuning

Shijia Yang*, Sheng Shen*, Tianjun Zhang*, Bohan Zhai, Joseph E. Gonzalez, Kurt Keutzer, Trevor Darrell

International Conference on Computer Vision (ICCV) 2023 Under Review

- Proposed the multitask vision-language prompt tuning (MVLPT) framework, including multitask prompt initialization and multitask prompt adaptation, and demonstrate the efficacy for each component
- Study the task transferability across 20 vision tasks with 400 combinations for each prompt tuning method to understand when MVLPT is most effective
- Evaluate MVLPT on the few-shot ELEVATER and cross-task generalization benchmarks, which sets the new state-of-the-art on 20-shot ELEVATER benchmark

INDUSTRY EXPERIENCE

Microsoft - DeepSpeed Team

Research Intern

January 2023 - Present

Bellevue, Washington

- Working on efficient inference of Large Language Models on long text, such as Llama2, MPT, and OPT
- Contributing LLM inference benchmarks to the open source DeepSpeed project

PERSONAL PROJECTS

Multi-modal Multi-task Learning for 3D Perceptive via Closed-loop Representation *advised by Prof. Yi Ma*

- Proposed a paradigm, inspired by closed-loop feedback from control theory, for a multi-modal multi-task 3D model
- Performed five common tasks for point-cloud perception, including foreground detection and object detection
- Achieved at least 4% mAP improvement for object detection on Argoverse dataset compared with UNet3D baseline

Snow Simulation *advised by Prof. Ren Ng*

- Used Taichi language designed for physical simulation and implemented the moving-least-square algorithm (MPM) to simulate realistic snow particles in 2D and 3D
- Studied snow particles' update rule for positions, affine velocity field, deformation gradient, etc.
- Implemented snowballs, frictional surface, and rigid body by adjusting particle distribution and parameters

Data Free Bert Quantization *advised by Prof. Sergey Levine*

- Alleviating the problem that quantized models may not have original data to perform post-quantization finetuning
- Proposed token generation algorithm to generate pseudo data for low precision Bert model quantization
- Improved accuracy of quantized Bert by 12.2%-61.7% on GLUE tasks compared to quantized Bert baseline